

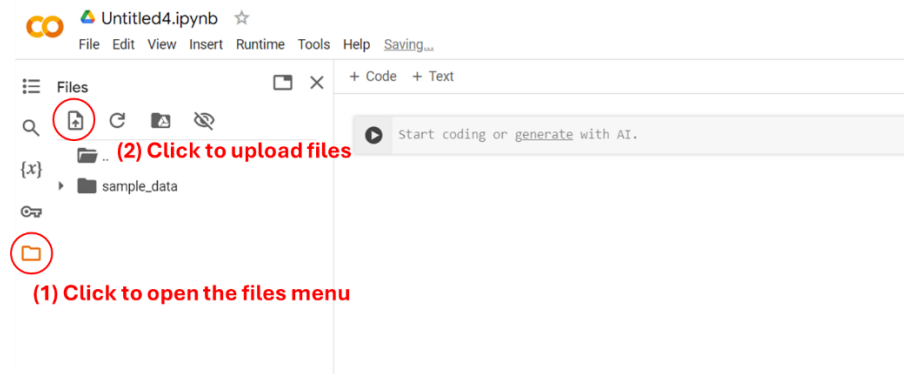
EE6310

Problem Set 11

Here, we will construct and train a neural network model through backpropagation. We will work with the logistic regression data set from the homework 8:

Metallic	H%	Li%	Be%	...	U%
y	x_1	x_2	x_3	...	x_{79}
1	0.41	0	0	...	0
0	0	0	0	...	0
0	0	0.04	0	...	0
...

The full data can be downloaded from the weekly learning menu named "HW11.csv." To upload the csv file to Google Colab, Follow the following steps:



First you can read the data set with the following command:

```
import pandas as pd
df = pd.read_csv('HW11.csv')
data = df.to_numpy()
Y = data[:,0]
X = data[:,1:]
```

We will utilize the functions that we have wrote in the previous homework. Run these codes to save the functions

```
def layer_sizes(X, Y):
    """
    Arguments:
    X -- input dataset of shape (number of examples, input size)
    Y -- labels of shape (number of examples, output size)
```

```

Returns:
n_0 -- the size of the input layer
n_1 -- the size of the hidden layer
n_2 -- the size of the output layer
"""

### START CODE HERE ### (≈ 3 lines of code)
n_0 = X.shape[1]
n_1 = 4
n_2 = 1
### END CODE HERE ###
return (n_0, n_1, n_2)

def initialize_parameters(n_0, n_1, n_2):
    """
    Argument:
    n_0 -- size of the input layer
    n_1 -- size of the hidden layer
    n_2 -- size of the output layer

    Returns:
    params -- python dictionary containing your parameters:
                W1 -- weight matrix of shape (n_0, n_1)
                b1 -- bias vector of shape (n_1)
                W2 -- weight matrix of shape (n_1, n_2)
                b2 -- bias vector of shape (n_2)
    """

    np.random.seed(2) # we set up a seed so that your output matches
    ours although the initialization is random.

    ### START CODE HERE ### (≈ 4 lines of code)

```

```

W1 = np.random.randn(n_0,n_1)*0.01
b1 = np.zeros((n_1,))
W2 = np.random.randn(n_1,n_2)*0.01
b2 = np.zeros((n_2,))
### END CODE HERE ###

parameters = {"W1": W1,
              "b1": b1,
              "W2": W2,
              "b2": b2}

return parameters

def forward_propagation(X, parameters):
    """
    Argument:
    X -- input data of size (m, n_0)
    parameters -- python dictionary containing your parameters
    (output of initialization function)

    Returns:
    A2 -- The sigmoid output of the second activation
    cache -- a dictionary containing "Z1", "A1", "Z2" and "A2"
    """
    # Retrieve each parameter from the dictionary "parameters"
    ### START CODE HERE ### (≈ 4 lines of code)
    W1 = parameters["W1"]
    b1 = parameters["b1"]
    W2 = parameters["W2"]
    b2 = parameters["b2"]
    ### END CODE HERE ###

```

```

# Implement Forward Propagation to calculate A2 (probabilities)
### START CODE HERE ### (≈ 4 lines of code)

Z1 = X@W1 + b1
A1 = np.tanh(Z1)
Z2 = A1@W2 + b2
A2 = 1/(1+np.exp(-Z2))
### END CODE HERE ###

cache = {"Z1": Z1,
        "A1": A1,
        "Z2": Z2,
        "A2": A2}

return A2, cache

```

- (a) Let's write the function to calculate the cost function by filling the code between "### START CODE HERE ###" and "### END CODE HERE ###"

Note that we have error in slides. The formula involving the calculation of cost function is:

$$L(\mathbf{A}^{[2]}, \mathbf{y}) = -\mathbf{y} \log(\mathbf{A}^{[2]}) - (1 - \mathbf{y}) \log(1 - \mathbf{A}^{[2]})$$

$$J(\mathbf{W}^{[1]}, \mathbf{b}^{[1]}, \mathbf{W}^{[2]}, \mathbf{b}^{[2]}) = \bar{L}$$

```

def compute_cost(A2, Y, parameters):
    """
    Computes the cross-entropy cost given in equation (13)

    Arguments:
    A2 -- The sigmoid output of the second activation, of shape
    (number of examples, 1)
    Y -- "true" labels vector of shape (number of examples)
    parameters -- python dictionary containing your parameters W1,
    b1, W2 and b2
    """

```

```
cost -- cross-entropy cost given equation (13)
"""

# Remove the feature (the last) dimension of the A2
### START CODE HERE ### (≈ 1 line of code)

### END CODE HERE ###

# Compute the cross-entropy cost
### START CODE HERE ### (≈ 2 lines of code)

### END CODE HERE ###

return Jprint ('The shape of X is: ' + str(shape_X))

n_1, n_2 = layer_sizes(X, Y)
parameters = initialize_parameters(n_0, n_1, n_2)
cache = forward_propagation(X, parameters)
compute_cost(A2, Y, parameters)

print("cost = " + str(J))

assert isinstance(J, float), 'Wrong answer!'

assert np.round(J,2) == 0.37
```

```
def backward_propagation(parameters, cache, X, Y):
    """
    Implement the backward propagation using the instructions above.
    """
```

Arguments:

parameters -- python dictionary containing our parameters
cache -- a dictionary containing "Z1", "A1", "Z2" and "A2".
X -- input data of shape (number of examples, n₀)
Y -- "true" labels vector of shape (number of examples)

Returns:

grads -- python dictionary containing your gradients with
respect to different parameters

```
"""
```

```
m = Y.shape[0]
```

```
# Retrieve also A1 and A2 from dictionary "cache".
```

```
### START CODE HERE ### (≈ 2 lines of code)
```

```
### END CODE HERE ###
```

```
# Add a dimension to Y
```

```
### START CODE HERE ### (≈ 1 lines of code)
```

```
### END CODE HERE ###
```

```
# Backward propagation: calculate dW1, db1, dW2, db2.
```

```
### START CODE HERE ### (≈ 6 lines of code, corresponding to 6  
equations on slide above)
```

```
### END CODE HERE ###
```

```
grads = {"dJdW1": dJdW1,  
         "dJdb1": dJdb1,  
         "dJdW2": dJdW2,
```

```

        "dJdb2": dJdb2}

    return grads

grads = backward_propagation(parameters, cache, X, Y)
print ("dJdW1 = "+ str(grads["dJdW1"]))
print ("dJdb1 = "+ str(grads["dJdb1"]))
print ("dJdW2 = "+ str(grads["dJdW2"]))
print ("dJdb2 = "+ str(grads["dJdb2"]))
assert grads["dJdW2"].shape == (4,1), 'Wrong answer!'
assert grads["dJdb2"].shape == (1,), 'Wrong answer!'
assert grads["dJdW1"].shape == (79,4), 'Wrong answer!'
assert grads["dJdb1"].shape == (4,), 'Wrong answer!'
assert np.round(np.mean(grads["dJdW1"]),5) == -0.00039, 'Wrong
answer!'

```

(c) Now let's use the calculated partial derivative to update the parameters:

```

def update_parameters(parameters, grads, learning_rate = 0.1):
    """
    Updates parameters using the gradient descent update rule given
    above

    Arguments:
    parameters -- python dictionary containing your parameters
    grads -- python dictionary containing your gradients

    Returns:
    parameters -- python dictionary containing your updated
    parameters
    """
    # Retrieve each parameter from the dictionary "parameters"

```

```

#### START CODE HERE #### (≈ 4 lines of code)


#### END CODE HERE ####

# Retrieve each gradient from the dictionary "grads"
#### START CODE HERE #### (≈ 4 lines of code)


## END CODE HERE ####

# Update rule for each parameter
#### START CODE HERE #### (≈ 4 lines of code)


#### END CODE HERE ####

parameters = {"W1": W1,
              "b1": b1,
              "W2": W2,
              "b2": b2}

return parameters

parameters = update_parameters(parameters, grads)

```



```

print("W1 = " + str(parameters["W1"]))
print("b1 = " + str(parameters["b1"]))
print("W2 = " + str(parameters["W2"]))
print("b2 = " + str(parameters["b2"]))

```

(d) Finally, let's write the function to train our model.

```

def nn_model(X, Y, n_1, num_iterations = 100000, print_cost=False):
    """
    Arguments:
    X -- dataset of shape (number of examples, n_0)
    Y -- labels of shape (number of examples)
    n_h -- size of the hidden layer
    num_iterations -- Number of iterations in gradient descent loop
    print_cost -- if True, print the cost every 10000 iterations

    Returns:
    parameters -- parameters learnt by the model. They can then be
    used to predict.
    """

    np.random.seed(3)
    n_0, _, n_2 = layer_sizes(X, Y)

    # Initialize parameters. Inputs: "n_0, n_1, n_2"
    ### START CODE HERE ### (≈ 1 lines of code)

    ### END CODE HERE ###

    # Loop (gradient descent)

```

```

    for i in range(0, num_iterations):

        ### START CODE HERE ### (≈ 4 lines of code)
        # Forward propagation. Inputs: "X, parameters". Outputs:
        "A2, cache".

        # Cost function. Inputs: "A2, Y, parameters". Outputs:
        "cost".

        # Backpropagation. Inputs: "parameters, cache, X, Y".
        Outputs: "grads".

        # Gradient descent parameter update. Inputs: "parameters,
        grads". Outputs: "parameters".

        ### END CODE HERE ###

        # Print the cost every 1000 iterations
        if print_cost and i % 10000 == 0:
            print ("Cost after iteration %i: %f" %(i, cost))

    return parameters

parameters = nn_model(X, Y, 4, num_iterations=100000,
print_cost=True)
print("W1 = " + str(parameters["W1"]))
print("b1 = " + str(parameters["b1"]))
print("W2 = " + str(parameters["W2"]))
print("b2 = " + str(parameters["b2"]))

```

(e) Let's write the code to make predictions using the trained model

```
def predict(parameters, X):
    """
    Using the learned parameters, predicts a class for each example
    in X

    Arguments:
    parameters -- python dictionary containing your parameters
    X -- input data of size (m, n_0)

    Returns
    y_hat -- vector of predictions of our model
    """

    # Computes probabilities using forward propagation, and
    # classifies to 0/1 using 0.5 as the threshold.
    ### START CODE HERE ### (≈ 2 lines of code)

    ### END CODE HERE ###

    # Remove the last dimension of the y_hat
    ### START CODE HERE ### (≈ 1 lines of code)

    ### END CODE HERE ###

    return y_hat
```

```

y_hat = predict(parameters, X)
print("predictions mean = " + str(np.mean(y_hat)))

assert np.round(np.mean(y_hat),3) == 0.545, 'Wrong answer!'
"A2": A2}

    return A2, cache

```

(f) Finally let's write the code to calculate accuracy. Google about how to calculate the accuracy for binary classification.

```

def caclulate_accuracy(y_hat, Y):
    """
    Using the y_hat and Y, calculate the accuracy

    Arguments:
    y_hat -- predicted Y values (m,)
    Y -- actual Y values (m, )

    Returns
    accuracy -- a float value of accuracy
    """

    #
    ### START CODE HERE ### (≈ 2 lines of code)
    accuracy = np.mean(Y==y_hat)
    ### END CODE HERE ###

    return accuracy

accuracy = caclulate_accuracy(y_hat, Y)

```

```
print("Accuracy = " + str(np.mean(y_hat)))  
assert np.round(np.mean(accuracy),3) == 0.837, 'Wrong answer!'
```